


# Chapter 3: Processes

*Presented By*  
**NARZU TARANNUM(NTR)**  
**DEPT. OF CSE, BRAC UNIVERSITY**





# Chapter 3: Topics Covered

- Process Concept
- Process Scheduling
- Operation on Processes
- Interprocess Communication






# Process Concept

- An operating system executes a variety of programs
- Textbook uses the terms *job* and *process* almost interchangeably



Batch System



Time Sharing  
System

- **Process** – a program in execution; process execution must progress in sequential fashion.
- A process includes:
  - program counter
  - stack
  - data section






# Multiple parts of a process

- The program code, also called **text section**
- Current activity including **program counter**, processor registers
- **Stack** containing temporary data, function parameters, return addresses, local variables, etc.
- **Data section** containing global variables
- **Heap** containing memory dynamically allocated during run time


- Process in Memory





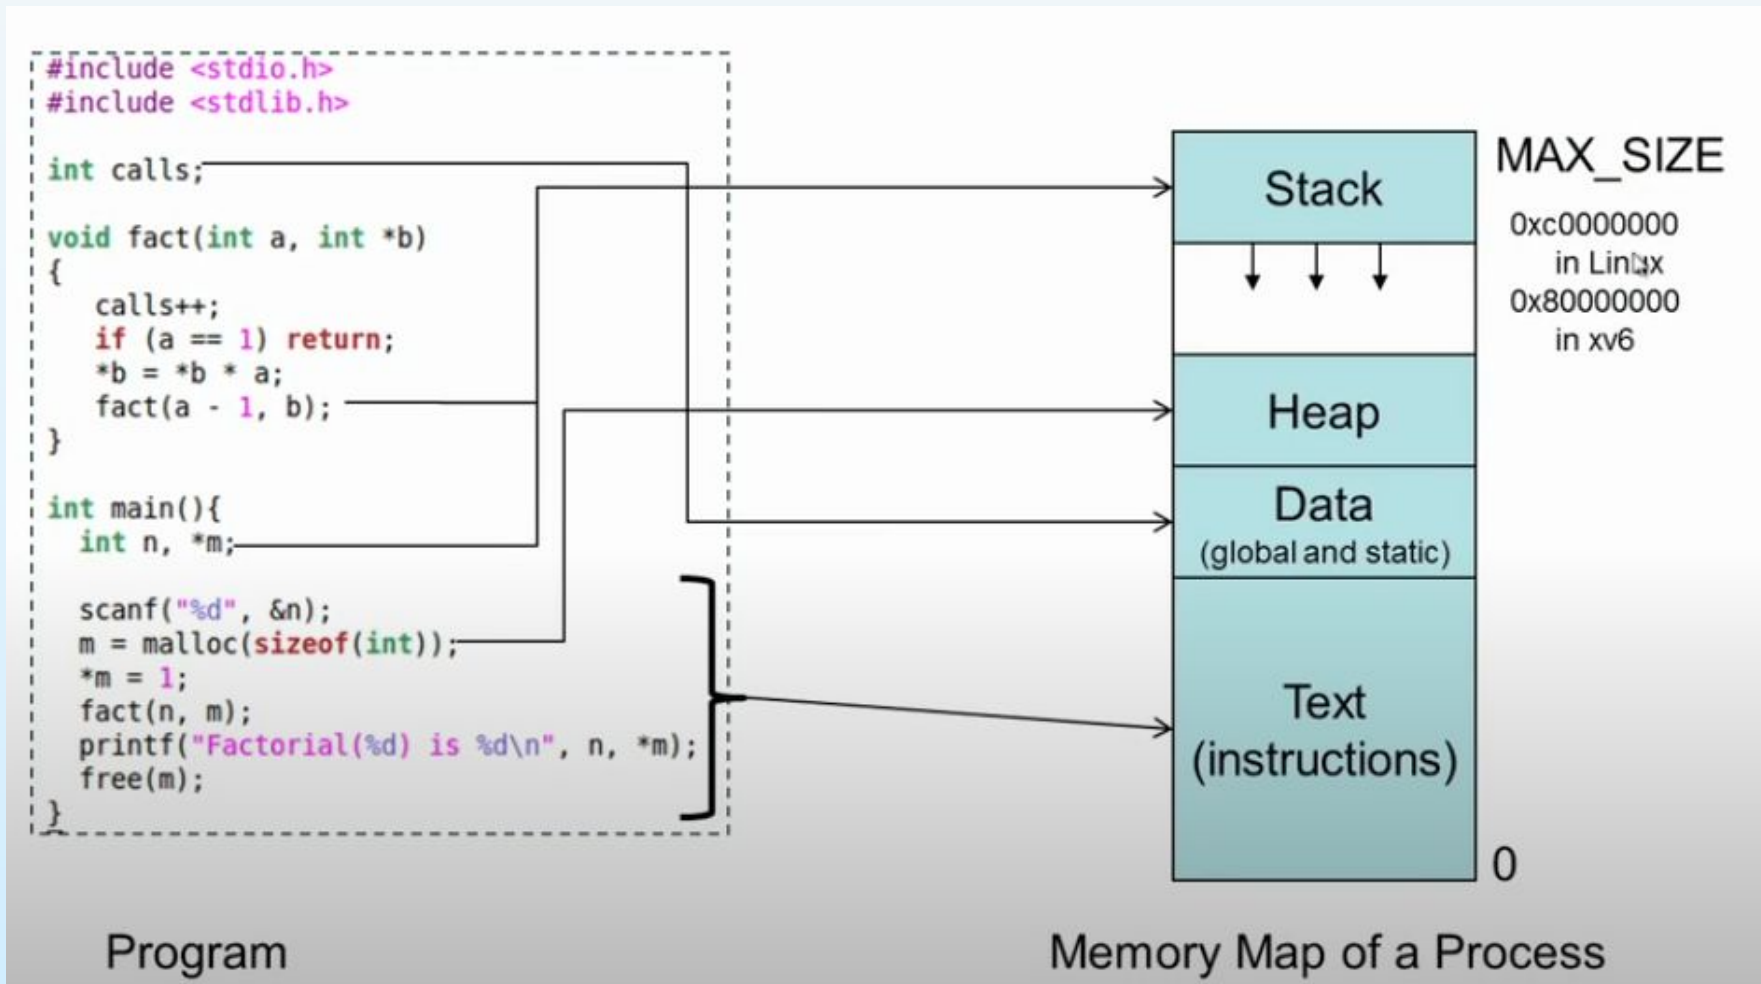
# Process Memory Map

```
#include <stdio.h>
#include <stdio.h>
Int calls;
Void fact(int a, int *b)
{
    calls++;
    if (a==1) return;
    *b = *b *a;
    fact(a-1,b);
}
Int main()
{
    int n, *m;
    scanf("%d", &n);
    m = malloc(sizeof(int));
    *m = 1;
    fact(n, m);
    print("Factorial (%d) is %d\n",n,*m);
    free (m);
}
```





# Process Memory Map





# Process Management

**A process is a program in execution.** It is a unit of work within the system.

Program is a *passive entity* stored on disk (**executable file**),

process is an *active entity*.

Process needs resources to accomplish its task

- CPU, memory, I/O, files

- Initialization data

Program becomes process when executable file loaded into memory.

Process termination requires reclaim of any reusable resources

Processes maybe single-threaded or multi-threaded

Process executes instructions sequentially, one at a time, until completion

Typically a system has many processes running concurrently on one or more CPUs

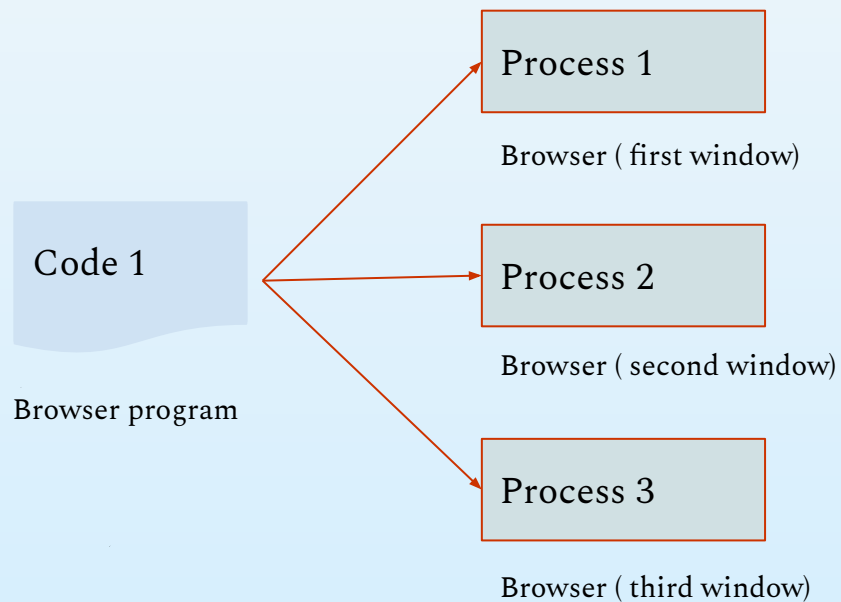
- Concurrency by multiplexing the CPUs among the processes / threads





# Same program, Different Process

- One program can be several processes
  - Consider multiple users executing the same program



Program code is same

Data, Heap, Stacks  
contains different  
information

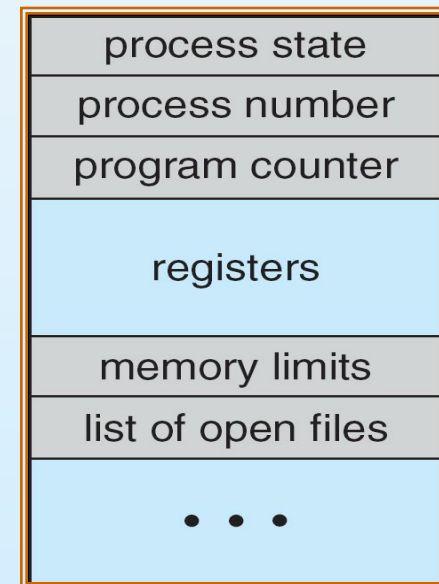






# Process Control Block (PCB)

- ❑ Information associated with each process is represented in the Operating system is Process Control Block(PCB).
- ❑ Process Control Block is a **data structure** that contains information of the process related to it.
  - ❑ Process id/number
  - ❑ Process state
  - ❑ Program counter
  - ❑ CPU registers
  - ❑ CPU scheduling information/ Priority
  - ❑ Memory-management information/ protection
  - ❑ List of open files
  - ❑ I/O status information





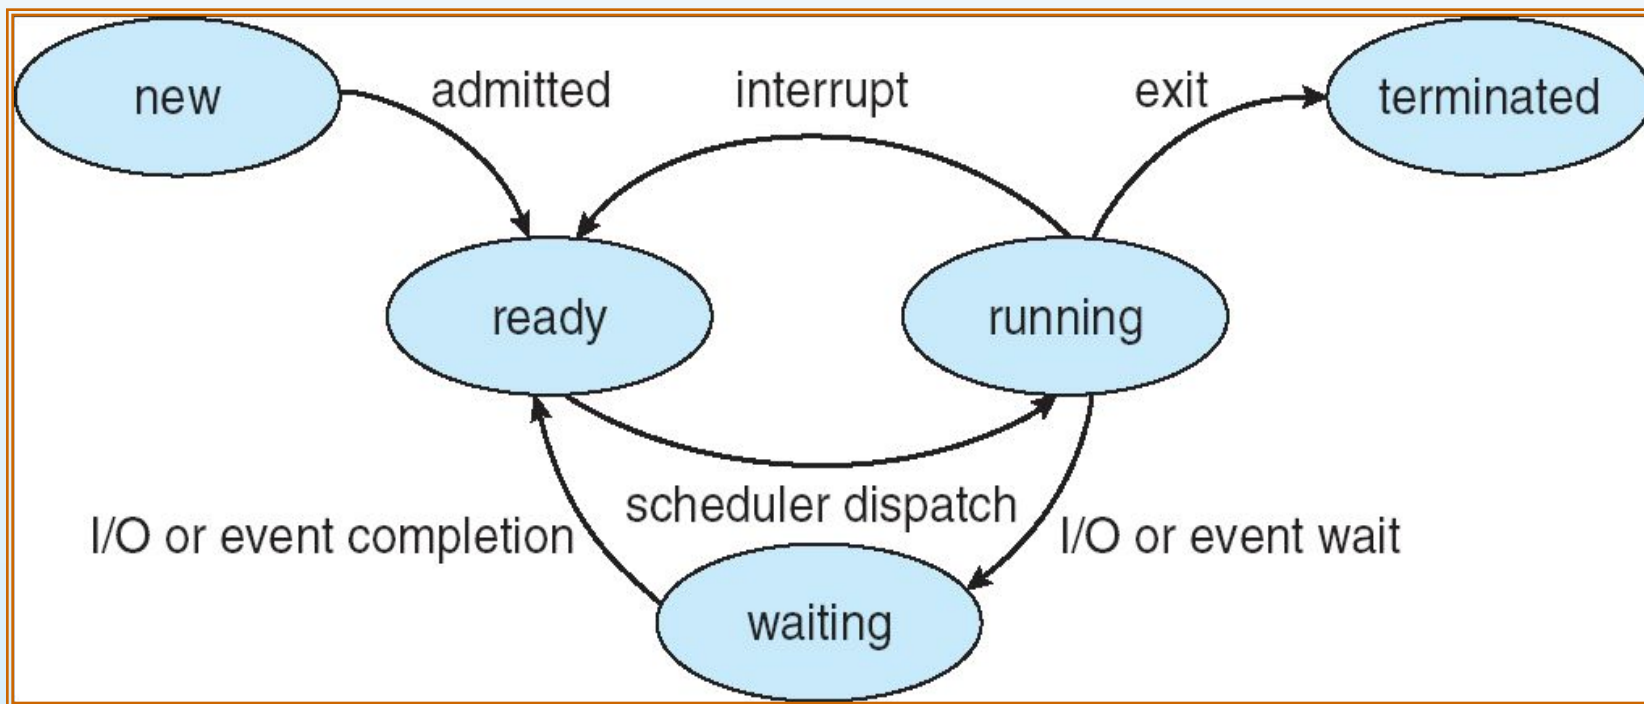
# Process State

- As a process executes, it changes *state*.
- *The state of a process is defined the current activity of that process.*
  - **new:** *The process is being created.*
  - **ready:** *The process is waiting to be assigned to a processor.*
  - **running:** *Instructions are being executed.*
  - **waiting:** *The process is waiting for some event to occur.*
  - **terminated:** *The process has finished execution.*



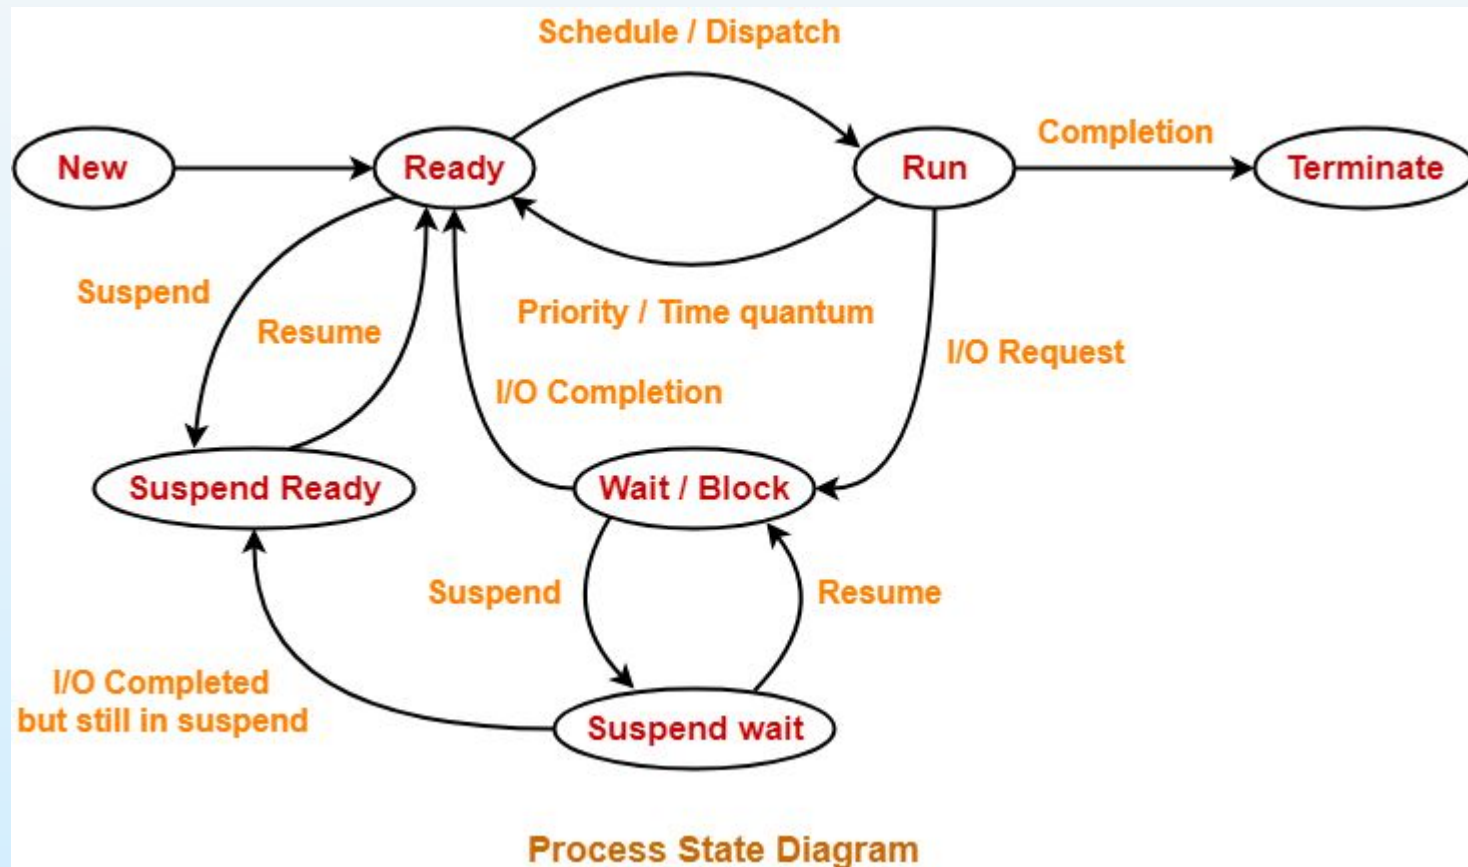


# Diagram of Process State





# Diagram of Process State





# Context Switch

- Context switching operations happen frequently on general purpose system.
- When CPU switches to another process, the system must **save the state of the old process and load the saved state for the new process**.
- Context-switch time is **overhead**; the system does no useful work while switching.
- **Context** of a process represented in the PCB
- The more complex the OS and the PCB □ the longer the context switch
- Time dependent on hardware support

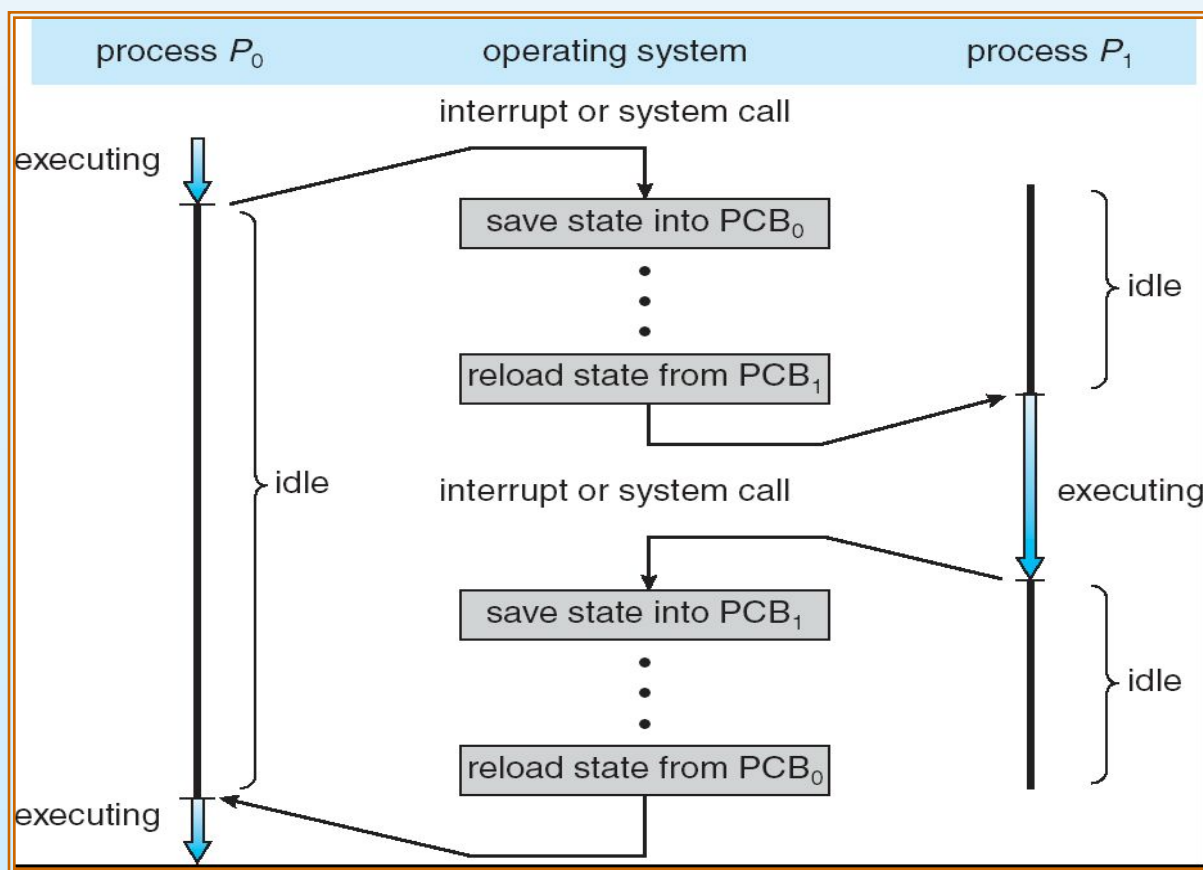




# CPU Switch From Process to Process

When an interrupt occurs, the system needs to save the current **context** (state) of the process running on the CPU.

- Context Switch:
1. Storing currently executed process context
  2. Restoring the next process context to execute






# Process Scheduling






# Process Scheduling

Multiple process is ready to execute.  
But, which Process should be executed first?



P1 -> P3 -> P2 -> P4 -> P5 ?  
P1 -> P2 -> P3 -> P5 -> P4 ?



CPU expecting processes to execute

Processes needs to be executed

- The objective of Multi-programming is to have some process running at all times, to maximize CPU utilization.
- The objective of Time sharing is to switch the CPU among processes so frequently that users can interact with each program while it is running.
- To meet these objectives, the process schedulers select an available process for program execution on CPU.







# Process Scheduling


- A process migrates among the various scheduling queues throughout its lifetime.
- Schedulers select processes from different queues to be passed to the next phase.
- The selection process is carried out by the appropriate scheduler.





# Schedulers

- ❑ **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
  - Long-term scheduler is invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)
  - The long-term scheduler controls the *degree of multiprogramming*
  - Long-term scheduler strives for good **process mix**
- ❑ **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
  - Short-term scheduler is invoked very frequently



Reside in Secondary Memory  
Keeps all the processes of the system

Reside in Main Memory  
Keeps all the processes that are waiting to be executed.





# Types of Process

- Processes can be described as either:
  - I/O-bound process** – spends more time doing I/O than computations.
  - CPU-bound process** – spends more time doing computations.


**Long Term Scheduler must select wisely !**

What will happen if all processes are I/O bound ?

=> Empty ready queue

What will happen if all processes are CPU bound ?

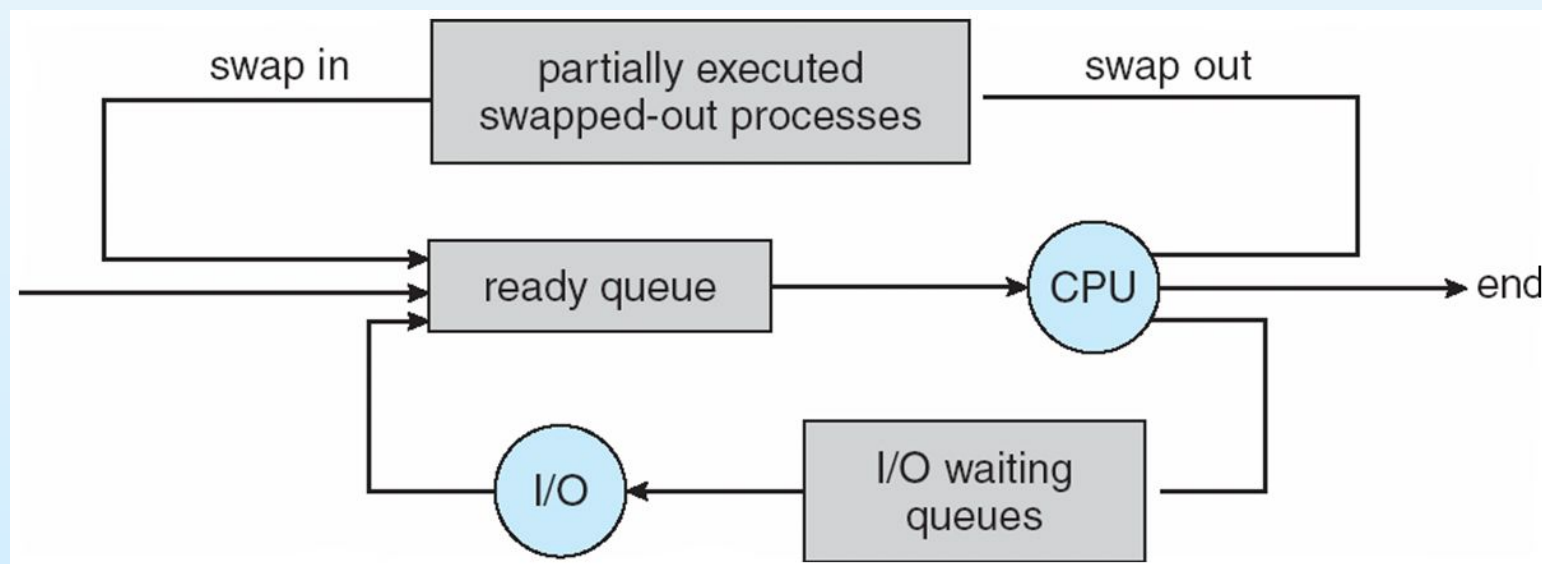
=> Empty waiting queue





# Mid-term scheduler

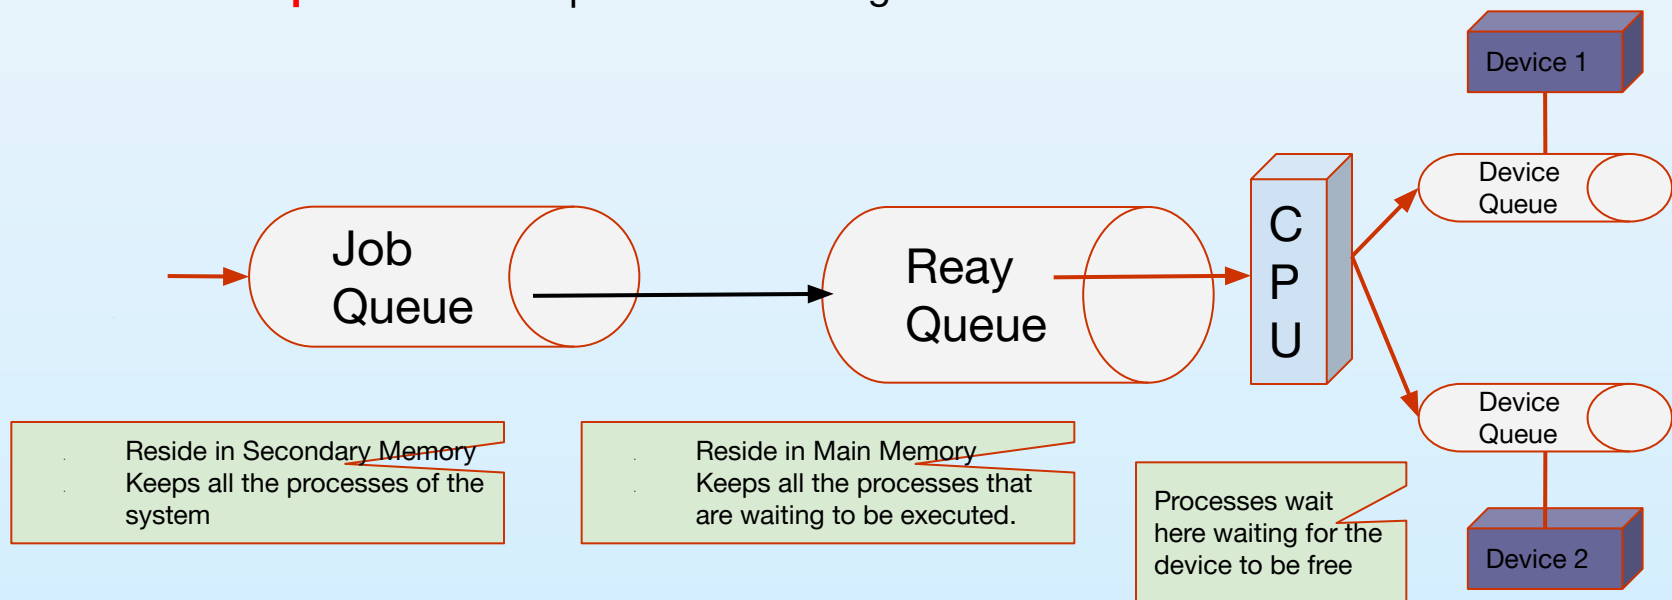
- ❑ The medium term scheduler swaps processes in and out memory to optimize CPU uses and manage memory allocation.
- ❑ sometime it can be advantageous to remove process from the memory and thus to reduce the degree of multi-programming.
- ❑ Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**. Swapping allows the system to pause and later resume a process, improving overall system efficiency.





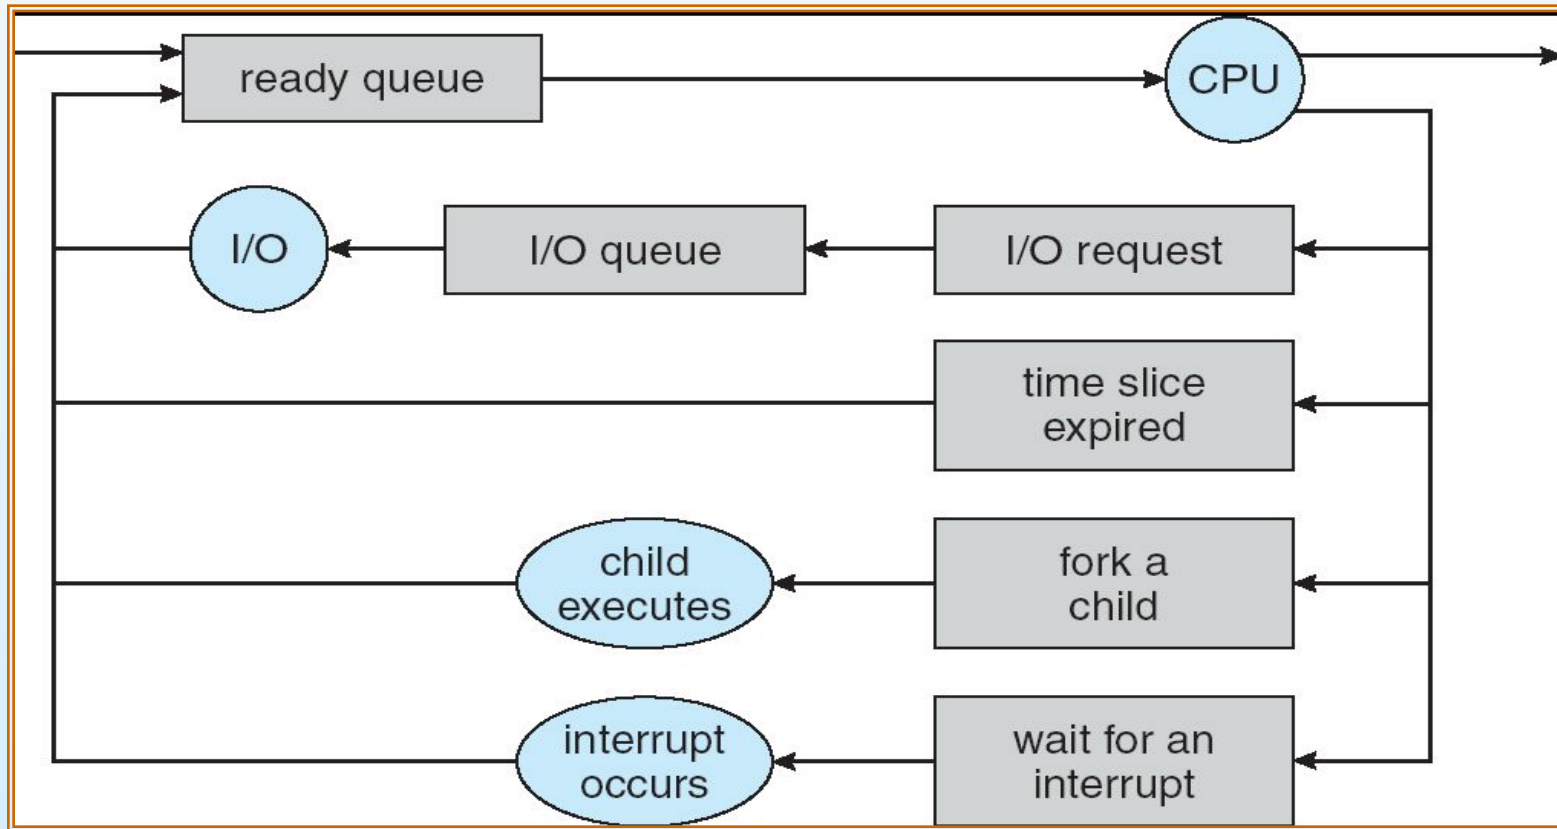
# Process Scheduling Queues

- ❑ Different queues are maintained for different state of the process.
- ❑ Processes migrate among the various queues
- ❑ Maintains scheduling queues of process:
  - **Job queue** – set of all processes in the system
  - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
  - **Device queues** – set of processes waiting for an I/O device





# Process Scheduling Queue/Queueing diagram





# Process scheduling as a queueing diagram

- A new process is initially put in the ready queue. It waits in the ready queue until it is selected for execution and is given the CPU. Once the process is allocated the CPU and is executing, one or several events could occur.
  - The process could issue an I/O request, and then be placed in an I/O queue.
  - The process could create a new sub process and waits for its termination.
  - The process could be removed forcibly from the CPU, as a result of an interrupt, and be put back to ready queue.
- A process continue this cycle until it terminates.





# Hardware Protection

The approach taken by most computer system is to provide hardware support that allows us to differentiate among various modes of execution.

**Dual-mode operation** allows OS to protect itself and other system components from errant users

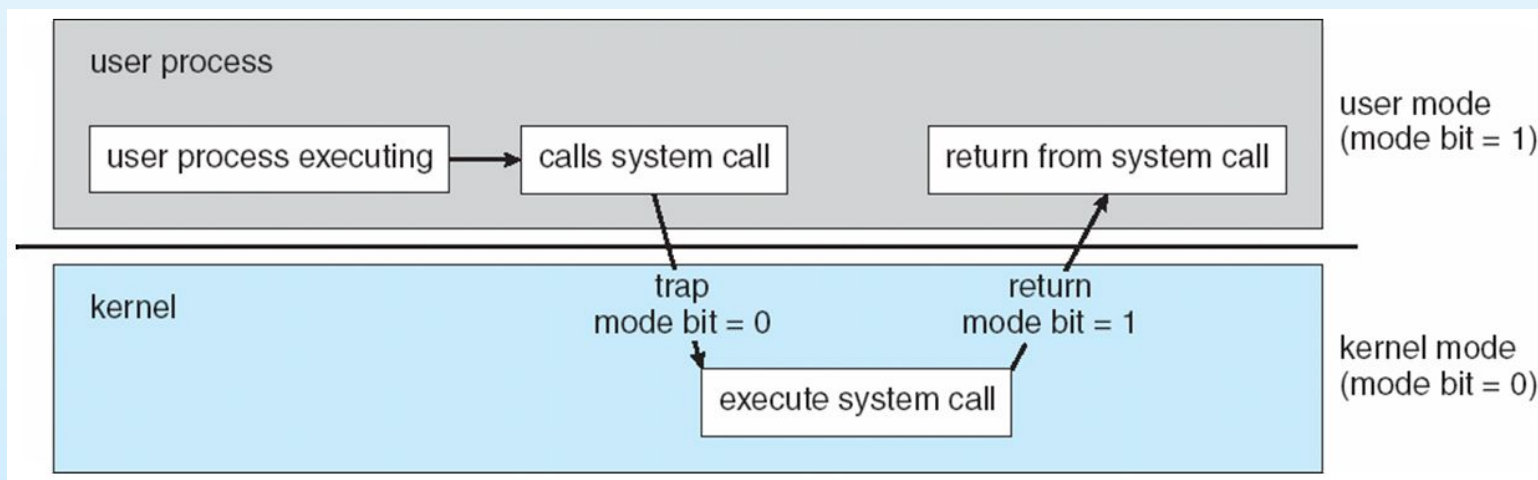
A bit, called the mode bit, is added to the hardware of the computer to indicate the current mode:

## User mode (1) and Kernel/Monitor/System mode (0)

Provides ability to distinguish when system is running user code or system code

Some instructions designated as **privileged**, only executable in system mode

System call changes mode to kernel, return from call resets it to user







# System Calls

- ❑ When a user application requests a service from the operating system (via a system call), the system must transition from user to kernel mode to fulfill the request. This is shown in previous figure.
- ❑ A **system call** is a way for programs to interact with the **operating system**.
- ❑ A computer program makes a **system call** when it makes a request to the **operating system's** kernel. **System call** provides the services of the **operating system** to the user programs via Application Program Interface(API).
- ❑ These calls are generally available as routines.





# System Call Implementation

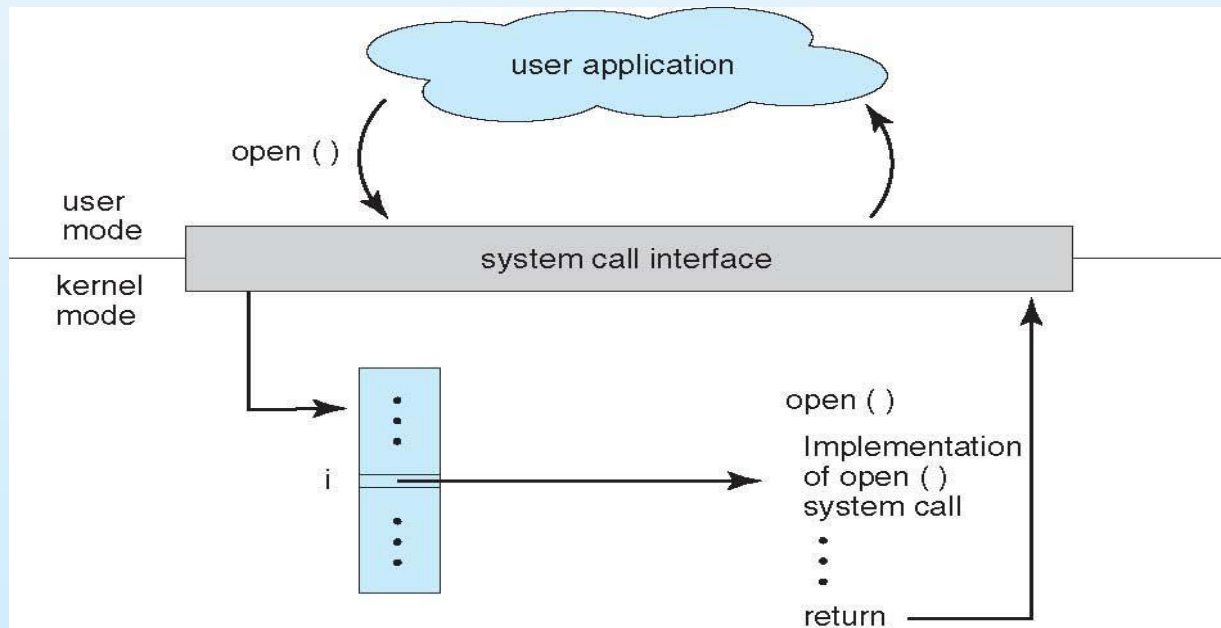
Typically, a number associated with each system call

**System-call interface** maintains a table indexed according to these numbers

The system call interface invokes the intended system call in OS kernel and returns status of the system call and any return values

The caller need know nothing about how the system call is implemented

API – System Call – OS Relationship





# Types of System Call

Type	Windows OS	Linux OS
Process Control	<b>CreateProcess() ExitProcess() WaitForSingleObject()</b>	<b>fork() exit() wait()</b>
File Manipulation	<b>CreateFile() ReadFile() WriteFile() CloseHandle()</b>	<b>open() read() write() close()</b>
Device Manipulation	<b>SetConsoleMode() ReadConsole() WriteConsole()</b>	<b>ioctl() read() write()</b>





# Types of System Call

Type	Windows OS	Linux OS
Information Maintenance	<b>GetCurrentProcessID() SetTimer() Sleep()</b>	<b>getpid() alarm() sleep()</b>
Communication	<b>CreatePipe() CreateFileMapping() MapViewOfFile()</b>	<b>pipe() shm_open() mmap()</b>
Protection	<b>SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()</b>	<b>chmod() umask() chown()</b>





# Operation on process

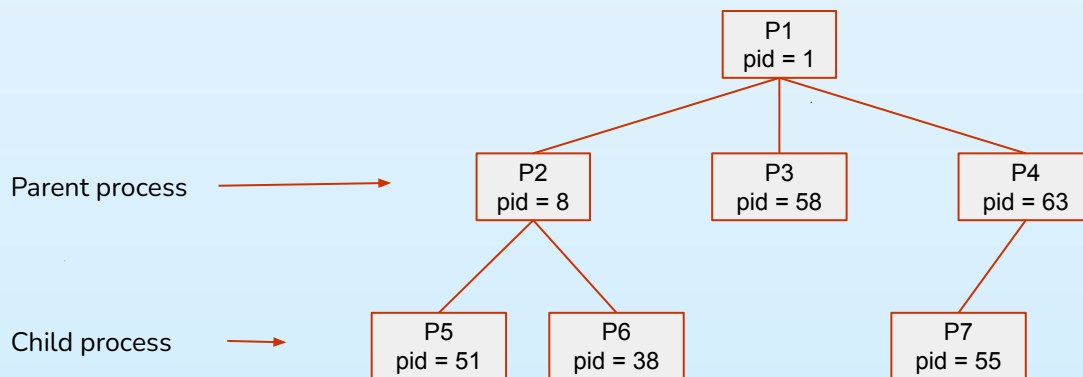
The processes in the system can execute concurrently, and must be **created** and **deleted** dynamically. OS must provide a mechanism for process creation and termination.





# Process Creation

- ❑ A process may create several new processes via a **create()** process system call during its execution.
- ❑ The creating process is called a **parent process**, and the new processes are called the **children** of that process
- ❑ Each of the new processes may create other processes, forming a tree of processes.
- ❑ Every process identified by a unique PID.
- ❑ Child process obtain resources from OS or are restricted to Parent's resources.
- ❑ Parent process may pass initializing data to child process





# Process Creation

When a process creates new process -

- The parent continues to execute concurrently with its children

Or,

- The parent waits until some or all of its children have terminated

Two address-space possibilities for the new process -

The child process is a duplicate of the parent process

Or,

The child process has a new program loaded into it.





# Process Creation in UNIX

- System Call: offers the services of the operating system to the user programs.
- **fork()**: create a new process, which becomes the child process of the caller
- **exec()**: runs an executable file , replacing the previous executable
- **wait()**: suspends execution of the current process until one of its children terminates.
- **exit()**: is used to terminate program execution.

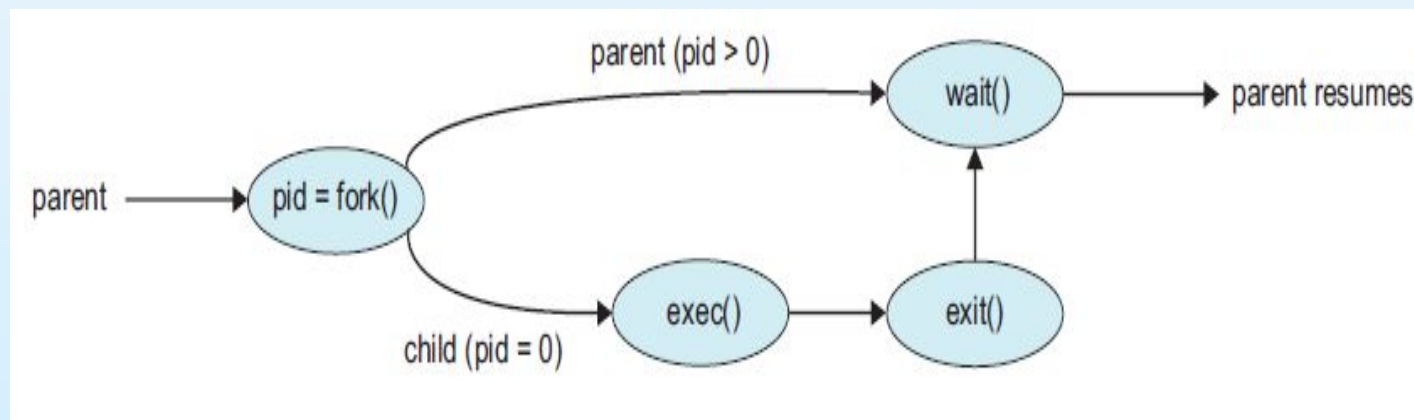


Fig: Process creation using fork() system call







# Process Creation in UNIX

- In Unix like operating system, fork system call is used for creating a new process where a process is a copy of itself, which is called the child process.
- The fork system call is usually a system call, implemented in the kernel.
- After a new child process is created, both processes will execute the next instruction following the fork() system call.
- If n times fork then total number of process created is  $2^n$  in which number of child process created is  $2^n - 1$

```
int main(){  
    fork();  
    fork();  
    printf("A");  
}
```

```
int main(){  
    fork();  
    fork();  
    fork();  
    printf("A");  
}
```





```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

**Figure 3.9** Creating a separate process using the UNIX `fork()` system call.





# Process Creation in UNIX (Example)

```
int main(){
    a = fork();
    if(a==0) fork();
    fork();
    printf("A");
}
```

```
int main(){
    int x = 1;
    a = fork();
    if(a==0){
        x = x -1;
        printf("value of x is: %d", x);
    }
    else if (a>0){
        wait(NULL);
        x = x +1;
        printf("value of x is: %d", x);
    }
}
```





# Process Creation in UNIX (Example)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid, pid1;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        pid1 = getpid();
        printf("child: pid = %d",pid); /* A */
        printf("child: pid1 = %d",pid1); /* B */
    }
    else { /* parent process */
        pid1 = getpid();
        printf("parent: pid = %d",pid); /* C */
        printf("parent: pid1 = %d",pid1); /* D */
        wait(NULL);
    }

    return 0;
}
```

**Figure 3.34** What are the pid values?





```
int main(){
    int id;
    static int x = 10;
    int y = 5;
    id = fork();
    if (id < 0){
        printf("fork failed\n");
    }
    else if(id == 0){
        printf("child started\n");
        printf("child finished\n");
    }
    else{
        wait(NULL);
        printf("parent started\n");
        x=x-2;
        y=y+5;
        printf("values of x: %d & y: %d\n",x,y);
        printf("parent finished\n");
    }
    x=x+5;
    y=y-5;
    printf("values of x: %d & y: %d\n",x,y);
    printf("terminating\n");

    return 0;
}
```





# Process Termination

- ❑ A process terminates when it finishes executing its last statement and asks the operating system to delete it using the `exit()` system call.
- ❑ At that point the process may return a status value (typically an integer) to its parent process (via the `wait()` system call).
- ❑ All the resources of the process—including physical and virtual memory, open files, and I/O buffers—are deallocated by the OS.





# Process Termination (Termination cause by another process)

- ❑ Parent may terminate the execution of children processes using the **abort ()** system call for a variety of reasons:
  - ❑ Child has exceeded allocated resources.
  - ❑ Task assigned to child is no longer required.
  - ❑ The parent is exiting and the operating systems does not allow a child to continue if its parent terminates ( cascading termination).

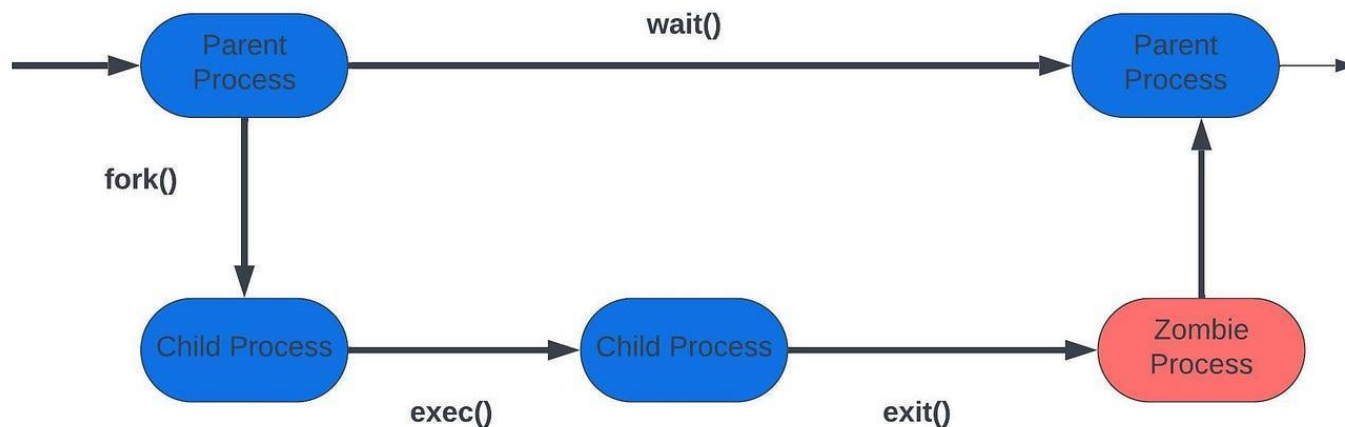




# Zombie Process in Unix

- A **zombie** process is a child process whose parent is yet to acknowledge its termination. On the flip side, an **orphaned** process is a child process whose parent has finished execution before it.

## Zombie Process







# Inter-process Communication

- Processes **within a system** may be ***independent*** or ***cooperating*** processes can execute concurrently.
- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process ( any process that shared data with other process is a cooperating process)
- **Advantages or reason of cooperating process**
  - 4 Information sharing
  - 4 Computation speed-up
  - 4 Modularity
  - 4 Convenience





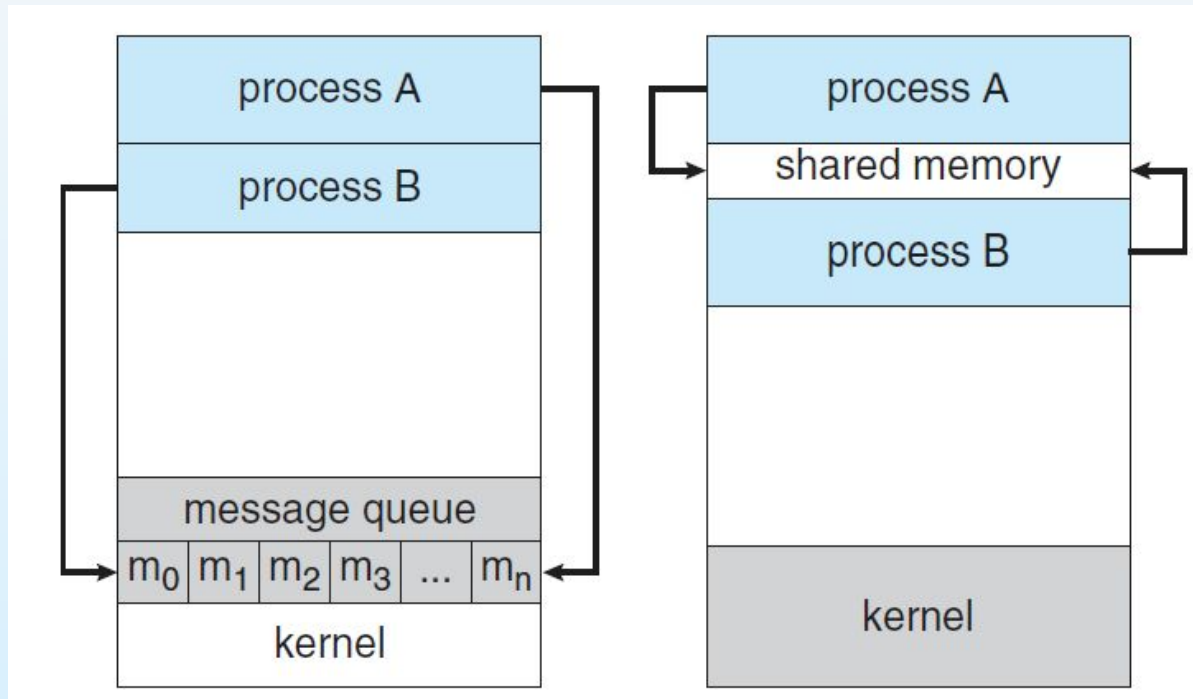
# Interprocess Communication

- Cooperating processes required **interprocess communication (IPC)** mechanism that will allow them to exchange data and information.
  - IPC provides a mechanism to allow processes to communicate and to synchronize there actions.
  - 4 There are two fundamental models of IPC
    - **Shared memory**
    - **Message passing**





# Interprocess Communication (IPC) Models



a) Message Passing

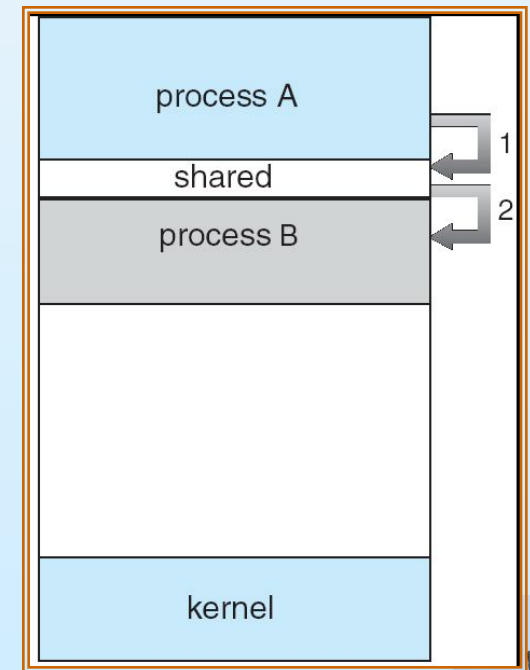
b) Shared Memory





# Shared-Memory


- Establish a region of shared memory for IIPC between processes.
- Other process that wish to communicate using this shared memory segment must attach it to their address space.
- Read and write data in the shared area
- Processes responsible for synchronization
- Must ensure that same memory location is not being modified by multiple processes at the same time.
- Shared memory requires that two or more processes agree to remove this restriction.





# Message-Passing systems

- In the message passing model, communication takes place by means of messages exchanged between the cooperating processes
- Message passing provides a mechanism to allow processes to communicate and to synchronize their actions without sharing the same address space.
- IPC facility provides at least two operations in message passing system:
  - **send**(message) – message size fixed or variable
  - **receive**(message)
- If  $P$  and  $Q$  wish to communicate, they need to:
  - establish a *communication link* between them
  - exchange messages via send/receive






# Shared Memory System

## (Producer-Consumer Problem)

Producer: produces products for consumer

Consumer: consumes products provided by producer





# Producer-Consumer Problem (Producer)

```
item next_produced;  
  
while (true) {  
    /* produce an item in next_produced */  
  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```


**in:** next free position in buffer  
**out:** first full position in buffer

Both initialized with 0.


in = 0  
out = 0

Here, BUFFER\_SIZE = 7

When buffer is full,  
in = 6, out = 0



When buffer is not full,  
In = 4, out = 0





# Producer-Consumer Problem (Consumer)

```
item next_consumed;


while (true) {
    while (in == out)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;


    /* consume the item in next_consumed */
}
```

Here, BUFFER\_SIZE = 7

When buffer is empty,  
in = 0, out = 0



When buffer is not empty,  
In = 5, out = 0





# End of Chapter 3

